

# **Architectural Paradigms of AI-First Web Standards: A Comprehensive Evaluation of the llms.txt and llms-full.txt Specifications**

The evolution of the global information ecosystem has reached a critical juncture where the primary consumers of web-based data are transitioning from human readers to autonomous artificial intelligence agents. For three decades, the foundational architecture of the World Wide Web has been optimized for visual rendering, utilizing HyperText Markup Language (HTML) to deliver rich, interactive experiences to human users. However, the emergence of Large Language Models (LLMs) and sophisticated Retrieval-Augmented Generation (RAG) frameworks has exposed the inherent inefficiencies of this human-centric design when processed by transformer-based architectures. The traditional web, characterized by a high ratio of boilerplate code to substantive information, presents a significant technical barrier to AI efficiency, primarily due to the finite nature of context windows and the computational cost of token processing.<sup>1</sup> In response to these challenges, the llms.txt and llms-full.txt standards have been proposed to serve as a machine-readable handshake between site owners and AI systems, providing a curated, high-signal guide to a domain's most critical knowledge.<sup>1</sup>

## **The Technical Divergence from Search-Centric Indexing to Reasoning-Centric Inference**

To understand the utility of the llms.txt specification, one must first delineate the fundamental differences between traditional search engine optimization (SEO) and the emerging field of generative engine optimization (GEO). Traditional SEO is built upon the requirement of indexing for discovery; search bots like Googlebot or Bingbot traverse a site to create a keyword-mapped index that users query to find relevant links.<sup>6</sup> The primary goal is traffic referral. Conversely, AI agents interact with content at the point of inference, where the objective is not merely to find a link but to extract, synthesize, and reason through information to provide a direct answer to a user.<sup>1</sup> This transition necessitates a shift in documentation strategy. While a robots.txt file provides a binary rulebook for access permissions, and a sitemap.xml offers a comprehensive inventory for discovery, neither provides the semantic guidance required for an LLM to navigate a complex site structure under strict token constraints.<sup>4</sup>

The standard proposed by Jeremy Howard and the team at Answer.AI identifies the context window as the most significant bottleneck in AI-web interactions.<sup>1</sup> An LLM attempting to parse a typical documentation page is frequently overwhelmed by navigation menus, sidebars,

advertisements, and tracking scripts that offer zero semantic value to the reasoning task.<sup>1</sup> This extraneous markup consumes tokens that would otherwise be allocated to the substantive documentation, leading to premature truncation of the context window or a dilution of the attention mechanism's focus.<sup>1</sup> The /llms.txt standard solves this problem by offering a plain-text, Markdown-based executive summary served at the root of the domain, designed to be consumed by agents at inference time.<sup>1</sup>

Standard Component	Description	Technical Requirement	Primary Beneficiary
Project Identity (H1)	The definitive name of the project or organization.	Required; must be the first line. <sup>1</sup>	LLM Entity Recognition
Project Summary	A concise blockquote summarizing the site's purpose.	Required; follows H1. <sup>1</sup>	Agentic Triage
Detailed Context	Non-heading paragraphs explaining site conventions.	Optional; provides nuanced usage tips. <sup>12</sup>	Contextual Alignment
Resource Lists (H2)	Thematic groupings of high-value URLs.	Recommended; uses Markdown link syntax. <sup>1</sup>	Targeted Retrieval
Optional Section	Resources that can be omitted for smaller context windows.	Specifically labeled H2 "Optional". <sup>9</sup>	Token Budgeting

## Comparative Utility of Curated Navigation and Full Corpus Ingestion

The dual-file approach, encompassing both /llms.txt and /llms-full.txt, represents a tiered strategy for context management. The /llms.txt file serves as a lightweight introduction, functioning as a "map" that lists the most important pages with brief, one-line descriptions.<sup>1</sup>

This allows an AI agent to perform a quick triage of the site's contents, identifying the specific sub-resources that are likely to contain the answer to a user's query.<sup>4</sup> For instance, a developer using an AI coding assistant may only need information regarding a specific API endpoint; the llms.txt file directs the agent to the relevant documentation page, bypassing the hundreds of other pages that might otherwise clutter the retrieval process.<sup>4</sup>

In contrast, the /llms-full.txt file provides the "full terrain" in a single, flattened Markdown document.<sup>4</sup> This comprehensive version is designed for agents with larger context windows or for complex tasks where the relationship between multiple documentation pages must be understood concurrently.<sup>3</sup> The utility of llms-full.txt is particularly pronounced in preventing the "fragmented retrieval" problem inherent in standard RAG systems.<sup>4</sup> When an agent relies on fragmented chunks of data, it often lacks the necessary "glue" or transitional context that connects disparate modules.<sup>4</sup> By offering the entire corpus in a single Markdown file, the site owner enables the LLM to perform "Context-Augmented Generation" (CAG), which eliminates the latency associated with multiple network calls for individual pages and ensures that the model can resolve cross-references with high fidelity.<sup>4</sup>

The economic implications of this efficiency are substantial. Token consumption is directly tied to API costs and inference latency. Research suggests that processing raw HTML can result in an 80% overhead compared to semantic Markdown.<sup>10</sup> By serving an llms-full.txt file, organizations can drastically reduce the token count required to provide a model with complete project context, thereby decreasing the total cost of ownership for AI-driven development tools and enhancing the user experience through faster response times.<sup>10</sup> This efficiency also mitigates the risk of hallucinations; when a model is provided with authoritative, clean text, it is significantly less likely to interpolate missing information from its training weights, leading to more reliable and verifiable outputs.<sup>4</sup>

## **Architectural Analysis: Implementation in Open-Source Serverless Environments**

The implementation of the llms.txt standard in a complex open-source project, such as the aws-cloudformation-media-downloader, illustrates the practical value of AI-first documentation for distributed cloud architectures.<sup>18</sup> This repository represents a sophisticated serverless application that integrates AWS Lambda, DynamoDB, S3, and API Gateway to manage media downloads via Feedly webhooks.<sup>18</sup> For an AI agent to assist a developer with this project, it must understand a multi-disciplinary stack that includes TypeScript runtime logic, OpenTofu (Terraform) infrastructure definitions, and third-party binaries like yt-dlp.<sup>18</sup>

A traditional crawl of such a repository would likely result in the AI agent ingesting thousands of lines of irrelevant configuration code, dependency lock files, and test fixtures before

identifying the core business logic.<sup>14</sup> By implementing an llms.txt file, the project owner can direct the agent specifically to the high-value documentation and source files. For example, the agent can be pointed toward the terraform/ directory to understand the Infrastructure-as-Code (IaC) layer and the layers/ directory to understand how external binaries are handled.<sup>18</sup> This targeted guidance is essential for complex features like the project's first-in-class ElectroDB adapter, which utilizes a single-table design in DynamoDB for optimized query performance.<sup>18</sup> Without a curated map, an AI might struggle to understand the relationship between the ElectroDB entities (Users, Sessions, Accounts) and the underlying AWS Global Secondary Indexes (GSIs).<sup>18</sup>

The project also includes a specialized convention capture system designed to preserve institutional knowledge.<sup>18</sup> This is a critical asset for AI agents, as it provides the stylistic and structural rules that the code must follow. By exposing these conventions through an llms.txt or llms-full.txt file, the developer ensures that the AI's suggestions are not merely syntactically correct but are also architecturally aligned with the project's philosophy.<sup>16</sup> The inclusion of API documentation generated from TypeSpec further enhances this utility, providing the AI with a machine-parseable contract for the backend services, which simplifies the generation of frontend integration code or testing suites.<sup>18</sup>

<b>Project Component</b>	<b>Documentation Requirement</b>	<b>Role of llms.txt</b>	<b>AI-Assisted Outcome</b>
AWS Infrastructure	Understanding resource dependencies (Lambda, SQS, SNS). <sup>18</sup>	Link to OpenTofu/Terraform templates. <sup>18</sup>	Accurate infrastructure modifications.
DynamoDB Schema	Explaining the ElectroDB single-table design. <sup>18</sup>	Pointer to entity and collection definitions. <sup>18</sup>	Efficient query generation.
Media Processing	Managing yt-dlp binary updates and FFmpeg conversion. <sup>18</sup>	Summary of binary download/verification logic. <sup>18</sup>	Robust processing pipeline maintenance.
API Specifications	Mapping TypeSpec to OpenAPI/Redoc	Direct link to structured API	Seamless service-to-service

	outputs. <sup>18</sup>	contracts. <sup>18</sup>	integration.
Build Pipeline	Defining pnpm security hardening and fixture sync. <sup>18</sup>	Guidance on build lifecycle and automation scripts. <sup>16</sup>	Secure and reproducible deployments.

## Strategy for Verification and Appropriate Usage by AI Agents

Ensuring that LLM agents utilize the provided llms.txt and llms-full.txt files appropriately requires a multi-faceted approach involving technical configuration, documentation structure, and active monitoring.<sup>8</sup> The foundational requirement is to host the files at the root of the domain (e.g., example.com/llms.txt) and serve them with the correct text/plain MIME type to ensure compatibility with standard HTTP clients used by LLMs.<sup>8</sup> Site owners must also verify that the files are publicly accessible and not obstructed by authentication barriers or aggressive firewall rules that might block AI scrapers.<sup>8</sup>

Once the technical hosting is secured, the next step is to test the files with various AI platforms to validate their semantic effectiveness.<sup>24</sup> This process, often referred to as "chatting with your docs," involves pasting the content of the llms.txt file into tools like ChatGPT, Claude, or Gemini and asking complex questions about the site's structure or specific features.<sup>12</sup> If the AI is unable to provide an accurate summary or fails to identify relevant links, it is an indication that the descriptions in the file are either too vague or use jargon that the model cannot resolve without additional context.<sup>8</sup> For technical projects like the aws-cloudformation-media-downloader, this might involve verifying that the AI understands the distinction between the "ElectroDB Adapter" and the "Better Auth" framework.<sup>18</sup>

Monitoring the usage of these files through server logs is the only way to confirm which agents are interacting with the documentation. Site owners should look for specific User-Agent strings such as GPTBot, ClaudeBot, and PerplexityBot.<sup>7</sup> Analyzing the frequency and sequence of requests can reveal if an agent is following the intended path—reading the /llms.txt index first before fetching specific sub-documents—or if it is defaulting to an inefficient broad crawl of the site.<sup>11</sup>

Verification Task	Method	Expected Result	Relevance for appropriate usage

Reachability Audit	curl -I https://yoursite.com /llms.txt	200 OK status; text/plain MIME type. <sup>21</sup>	Confirms the file is available for AI crawlers.
Semantic Validation	Manual prompting in Claude/ChatGPT. <sup>24</sup>	Accurate project summary and link identification. <sup>12</sup>	Ensures the model can "reason" using the file.
Bot Interaction Log	Grep server logs for AI User-Agents. <sup>7</sup>	Sequential access patterns from /llms.txt to specific URLs. <sup>22</sup>	Validates that agents are using the "map" correctly.
Tokenization Test	Compare HTML vs. Markdown token counts. <sup>10</sup>	Significant reduction in token usage for same content. <sup>10</sup>	Confirms the file provides an efficient context.
Directive Compliance	Audit AI responses for source attribution. <sup>29</sup>	Correct citation of the site as the source of truth. <sup>30</sup>	Ensures the brand is properly represented.

## The Role of Model Context Protocol (MCP) in the Agentic Workflow

The Model Context Protocol (MCP) has emerged as a critical transport layer for the llms.txt ecosystem. Introduced by Anthropic, MCP provides a standardized way for AI applications to interface with external data sources and tools without writing custom integration code for every API.<sup>14</sup> In the context of the aws-cloudformation-media-downloader project, an MCP server could be implemented to dynamically fetch the latest llms.txt and llms-full.txt files, parsing them to provide the LLM with real-time documentation during a development session.<sup>15</sup>

This protocol effectively acts as a "USB-C" for AI context, allowing any MCP-compliant agent to "plug into" a repository's documentation and instantly understand its structure.<sup>14</sup> For a developer, this means their AI assistant no longer relies on potentially outdated training data; instead, it uses the MCP server to retrieve the most recent TypeSpec definitions or Terraform changes directly from the root documentation files.<sup>14</sup> This setup is particularly useful for managing the "800k token problem," where an agent might otherwise exhaust its context window simply trying to understand the directory layout of a large project.<sup>14</sup> By using MCP tools like fetch\_llms\_txt, the agent can selectively ingest only the most relevant sections of the

documentation, maintaining a lean and focused context for the task at hand.<sup>14</sup>

The integration of MCP servers also provides a centralized mechanism for managing permissions and rate limiting. A public MCP server, such as the mcp-llmstxt server deployed on Cloudflare Workers, can help AI agents discover known implementations across multiple domains while ensuring that the host application is not overwhelmed by excessive scraper traffic.<sup>34</sup> This creates a more sustainable and governed ecosystem for AI interactions compared to unmanaged web scraping.<sup>33</sup>

## **Generative Engine Optimization (GEO) and the Business Value Proposition**

Beyond the technical benefits for developers, the adoption of llms.txt and llms-full.txt represents a strategic move in the realm of Generative Engine Optimization (GEO). As AI-powered search engines like Perplexity and ChatGPT Search begin to dominate the discovery phase, businesses must ensure that their content is not only accessible but also correctly interpreted by these models.<sup>2</sup> An accurately configured llms.txt file serves as a brand-controlled narrative, guiding the LLM toward official product descriptions, pricing policies, and support documentation.<sup>7</sup>

For a software project, this translates to improved user onboarding and a reduction in support tickets. When a developer asks an AI, "How do I deploy the media downloader to AWS?", the model will parse the llms.txt file, find the deployment guide, and provide an accurate set of steps based on the project's actual Terraform templates.<sup>6</sup> This "zero-click" success encourages adoption and builds trust in the project's documentation.<sup>4</sup> Furthermore, the presence of these files acts as a signal of technical sophistication and "AI-readiness," which is increasingly important for developer tools and open-source libraries looking to attract users in an AI-first market.<sup>6</sup>

The competitive advantage of early adoption is also worth noting. Organizations that implement these standards today are positioning themselves to be the "source of truth" in the training and retrieval datasets of tomorrow.<sup>7</sup> As the web becomes increasingly saturated with AI-generated noise, authoritative, machine-readable documentation served directly from the domain root will become the most reliable anchor for AI agents.<sup>13</sup> This proactive stance on content governance allows site owners to define their own parameters for attribution and usage, rather than having those terms dictated by the scraping policies of third-party AI companies.<sup>7</sup>

## **Governance, Consent, and the Future of AI Crawling Protocols**

The emergence of llms.txt is occurring alongside a broader debate regarding AI data sovereignty and the limitations of the thirty-year-old robots.txt standard.<sup>7</sup> Robots.txt was designed for a world where crawlers were merely creating a list of links; it was not built to govern the fine-tuning of multi-billion parameter neural networks on proprietary content.<sup>7</sup> While llms.txt is primarily a tool for inference and discovery, it also provides a framework for communicating training preferences.<sup>7</sup>

Site owners can use the llms.txt file to specify distinct rules for inference versus training. For example, a project may allow an AI to use its documentation for real-time question answering (inference) while prohibiting the same content from being used to train a model that might eventually compete with the project's core value proposition.<sup>7</sup> By including lines such as Disallow: /training or Attribution: Required, the site owner establishes a clear set of expectations for ethical AI behavior.<sup>30</sup> While these signals are not currently legally binding, they provide a technical foundation for future legal frameworks and industry-wide consent standards.<sup>30</sup>

The future of these standards likely involves deeper integration with browser-level controls and search console platforms. Google's recent addition of an llms.txt file to its own developer documentation portal, despite previous public skepticism, suggests that even the largest players in the search industry are preparing for a world where AI-specific crawling protocols are the norm.<sup>39</sup> As more organizations follow suit, we can expect to see a stabilization of the specification and the emergence of more sophisticated tools for validating, monitoring, and monetizing AI access to web-based knowledge.<sup>1</sup>

## **Practical Implementation Roadmap for Complex Codebases**

For the aws-cloudformation-media-downloader project, a comprehensive implementation of these standards should follow a structured approach that emphasizes architectural clarity and token efficiency. The primary objective is to enable an AI agent to assist with three core tasks: infrastructure deployment, codebase refactoring, and API integration.<sup>18</sup>

### **Step 1: Curating the Inference Index (/llms.txt)**

The root /llms.txt file should serve as the executive summary of the project. It should lead with a clear H1 title and a blockquote that identifies the project as a serverless AWS media downloader.<sup>1</sup> The subsequent H2 sections should group the links into logical categories such as ## Core Infrastructure, ## Application Logic, and ## API Reference.<sup>1</sup> Crucially, the descriptions for each link should be factual and concise, explaining exactly what information the model will find at that URL (e.g., "(terraform/): OpenTofu definitions for Lambda, S3, and DynamoDB resources").<sup>9</sup>

## Step 2: Generating the Comprehensive Corpus (/llms-full.txt)

The llms-full.txt file should be generated through an automated build process that flattens the project's documentation into a single Markdown document.<sup>1</sup> This file should include the full text of the README, the TypeSpec API definitions, the architectural conventions document, and the core TypeScript entity definitions from the ElectroDB layer.<sup>18</sup> By including the code for the ElectroDB entities, the project owner allows the LLM to understand the data schema without needing to browse the file system, which is essential for generating accurate DynamoDB queries.<sup>16</sup>

## Step 3: Providing Clean Markdown Equivalents (.md)

To further reduce token overhead, the project should offer clean Markdown versions of any documentation that is normally served as HTML.<sup>1</sup> For the Redoc-based API documentation, the project should ensure that the raw Markdown or OpenAPI JSON is accessible at a predictable URL (e.g., docs/api.md).<sup>1</sup> This allows the AI agent to bypass the HTML structure of the documentation viewer entirely, focusing only on the function signatures and data types.<sup>5</sup>

## Step 4: Configuring Governance and Monitoring

Finally, the project must align its llms.txt directives with its robots.txt file. If the goal is to encourage AI assistance for developers, the robots.txt should explicitly allow GPTBot and ClaudeBot access to the root directory and the documentation paths.<sup>7</sup> The llms.txt file should include a level of "meta-instruction" for the agent, such as ### Review Options that guide the AI on how to perform a code review based on the project's specific security hardening and supply chain protection scripts.<sup>1</sup> Regular monitoring of the server logs will then allow the developer to confirm that the AI agents are successfully discovering the files and using them to provide better, more context-aware suggestions.<sup>22</sup>

Implementation Milestone	Project-Specific Detail	Technical Action	AI Benefit
<b>Root Initialization</b>	aws-cloudformatio n-media-downloa der	Create /llms.txt with H1 and blockquote. <sup>1</sup>	Immediate site identification.
<b>Infrastructure Mapping</b>	OpenTofu/Terrafor m modules. <sup>18</sup>	Link terraform/ as a core documentation section. <sup>12</sup>	Improved IaC assistance.

<b>Entity Exposure</b>	ElectroDB single-table entities. <sup>18</sup>	Include full entity code in llms-full.txt. <sup>4</sup>	Reduced schema hallucinations.
<b>Binary Context</b>	yt-dlp update & check logic. <sup>18</sup>	Summarize the binary verification workflow. <sup>1</sup>	Better maintenance support.
<b>Convention Capture</b>	Institutional knowledge & style. <sup>18</sup>	Link to docs/conventions. md. <sup>18</sup>	On-brand code generation.
<b>Access Verification</b>	User-agent tracking. <sup>22</sup>	Verify GPTBot and ClaudeBot access in logs. <sup>7</sup>	Confirmed agentic usage.

In summation, the llms.txt and llms-full.txt standards offer a robust and technically sound solution to the challenges of AI-web interaction. For complex, cloud-native projects like the aws-cloudformation-media-downloader, these files are not merely an optional addition but are a critical component of the modern developer experience.<sup>5</sup> By optimizing for token efficiency, semantic clarity, and agentic discovery, project owners can ensure that their software is accurately understood and effectively supported by the rapidly growing ecosystem of AI tools and agents.<sup>1</sup> The transition to AI-first documentation is an inevitability of the agentic era, and early adoption remains the most effective strategy for maintaining brand authority and technical excellence in an increasingly automated web.<sup>1</sup>

## Works cited

1. LLMs.txt Explained | TDS Archive - Medium, accessed December 21, 2025, <https://medium.com/data-science/llms-txt-414d5121bcb3>
2. llms.txt: The New Frontier of AI Crawling and SEO - Xfunnel.ai, accessed December 21, 2025, <https://www.xfunnel.ai/blog/understanding-llms-2025>
3. Making ML Documentation AI-Friendly: ZenML's Implementation of llms.txt, accessed December 21, 2025, <https://www.zenml.io/blog/llms-txt>
4. What Is LLM.txt & Why It Matters for AI SEO - Wellows, accessed December 21, 2025, <https://wellows.com/blog/what-are-llm-txt/>
5. What is llms.txt? Breaking down the skepticism - Mintlify, accessed December 21, 2025, <https://www.mintlify.com/blog/what-is-llms-txt>
6. What Is LLM.txt (aka llms.txt), and Should You Use It? // - Singularity Digital Marketing, accessed December 21, 2025, <https://singularity.digital/insights/what-is-llms-txt/>
7. What Is LLMs.txt? & Do You Need One? - Neil Patel, accessed December 21, 2025, <https://neilpatel.com/blog/llms-txt-files-for-seo/>
8. What is LLMs.txt and How Can It Boost Your GEO Efforts? - Firebrand

- Communications, accessed December 21, 2025,  
<https://www.firebrand.marketing/2025/08/what-is-llms.txt-and-how-it-can-boost-your-geo-efforts/>
- 9. What Is Llms.txt? Reality vs. Hype | Similarweb, accessed December 21, 2025,  
<https://www.similarweb.com/blog/marketing/geo/llms-txt/>
  - 10. Debunking LLMs.txt Myths: What You Need to Know for AI Visibility - Wix.com, accessed December 21, 2025,  
<https://www.wix.com/studio/ai-search-lab/llms-txt-myths>
  - 11. How often do LLMs visit llms.txt? - Mintlify, accessed December 21, 2025,  
<https://www.mintlify.com/blog/how-often-do-llms-visit-llms-txt>
  - 12. llms-txt: The /llms.txt file, accessed December 21, 2025, <https://llmstxt.org/>
  - 13. LLMs.txt: The Emerging Standard Reshaping AI-First Content Strategy | ScaleMath, accessed December 21, 2025, <https://scalemath.com/blog/llms-txt/>
  - 14. The LLM.txt Directory MCP Server: Your AI's Guide to Codebases, accessed December 21, 2025,  
<https://skywork.ai/skypage/en/llm-directory-ai-codebases/1978668209129771008>
  - 15. Give Your AI Agents Deep Understanding With LLMS.txt | by Dazbo (Darren Lester) | Google Cloud - Medium, accessed December 21, 2025,  
<https://medium.com/google-cloud/give-your-ai-agents-deep-understanding-with-llms-txt-4f948590332b>
  - 16. Writing Docs for AI: Making Your Product Seamless for Cursor, Windsurf, and Claude Code Users - Rivet, accessed December 21, 2025,  
<https://www.rivet.dev/blog/2025-03-15-writing-docs-for-ai/>
  - 17. Working with llms.txt | Platform Overview - Mastercard Developers, accessed December 21, 2025,  
<https://developer.mastercard.com/platform/documentation/agent-toolkit/working-with-llmstxt/>
  - 18. jOnathan-II0yd/aws-cloudformation-media-downloader: A Terraform project for downloading media (e.g. YouTube videos) via Feedly - GitHub, accessed December 21, 2025,  
<https://github.com/jOnathan-II0yd/aws-cloudformation-media-downloader>
  - 19. aws-samples/packaged-vod-downloader: The Packaged VOD Downloader Workflow simplifies the process of repackaging HLS/TS Video On Demand assets to one or more output packaging formats (including HLS, DASH and CMAF). - GitHub, accessed December 21, 2025,  
<https://github.com/aws-samples/packaged-vod-downloader>
  - 20. Best Practices for Context Management when Generating Code with AI Agents, accessed December 21, 2025,  
<https://docs.digitalocean.com/products/gradient-ai-platform/concepts/context-management/>
  - 21. What is llms.txt? Why it's important and how to create it for your docs – GitBook Blog, accessed December 21, 2025,  
<https://www.gitbook.com/blog/what-is-llms-txt>
  - 22. What are AI user agents?, accessed December 21, 2025,  
<https://usehall.com/guides/what-are-ai-user-agents>

23. LLMS.txt Best Practices & Implementation Guide - Rankability, accessed December 21, 2025, <https://www.rankability.com/guides/llms-txt-best-practices/>
24. llms.txt & .md files - Important AI Visibility helper or hoax?, accessed December 21, 2025, <https://peec.ai/blog/llms-txt-md-files-important-ai-visibility-helper-or-hoax>
25. llms.txt for Nuxt Sites, accessed December 21, 2025, <https://nuxtseo.com/learn-seo/nuxt/controlling-crawlers/llms-txt>
26. How to Create llms.txt? 5 Points Following Dr. Howard Standard - Bursa Web, accessed December 21, 2025, <https://bursaweb.com/en/how-to-create-an-llms-txt-file-5-critical-things-you-must-know/>
27. The Ultimate Guide to llms.txt for GEO/AEO: What Every Business Leader Must Know, accessed December 21, 2025, <https://insidea.com/blog/seo/geo/llms-txt-for-geo-aeo/>
28. llms.txt for E-commerce: A Practical Guide to Preparing Your Site for AI Crawlers, accessed December 21, 2025, <https://www.tngshopper.com/post/llms-txt-for-e-commerce-a-practical-guide-to-preparing-your-site-for-ai-crawlers>
29. LLMs.txt vs LLM-Full.txt: What SEO Professionals Need to Know in 2025 | HackerNoon, accessed December 21, 2025, <https://hackernoon.com/llmstxt-vs-llm-fulltxt-what-seo-professionals-need-to-know-in-2025>
30. LLMs.txt File - Level Agency, accessed December 21, 2025, <https://www.level.agency/ai-seo-glossary/llms-txt-file/>
31. Future-Ready Content: Why Your Website Needs an llms.txt File - Matter. Design, accessed December 21, 2025, <https://matterdesign.com.au/news/future-ready-content-why-your-website-needs-an-llms-txt-file/>
32. Model Context Protocol - GitHub, accessed December 21, 2025, <https://github.com/modelcontextprotocol>
33. Model Context Protocol (MCP): A comprehensive introduction for developers - Stytch, accessed December 21, 2025, <https://stytch.com/blog/model-context-protocol-introduction/>
34. MCP llms.txt Server - LobeHub, accessed December 21, 2025, <https://lobehub.com/mcp/marthakelly-mcp-llmstxt>
35. LLMs.txt for Ecommerce (Get Products Discovered in AI Search) - BigCommerce, accessed December 21, 2025, <https://www.bigcommerce.com/blog/ecommerce-llms-txt/>
36. How to Get Your Website Found by AI in 2025 Using the llms.txt File - WolfPack Advising, accessed December 21, 2025, <https://wolfpackadvising.com/blog/boost-ai-seo-2025-with-llms-txt-file/>
37. The Ultimate LLM.txt Guide for Marketers and SEOs - Writesonic, accessed December 21, 2025, <https://writesonic.com/blog/llm-txt-guide>
38. LLMS.Txt: Everything You Need to Know - Open Forge AI, accessed December 21, 2025, <https://www.openforge.ai/blog/LLMS-Txt>

39. Google Adds LLMs.txt to Search Developer Docs Portal - Stan Ventures, accessed December 21, 2025,  
<https://www.stanventures.com/news/google-adds-llms-txt-to-search-developer-docs-portal-6066/>
40. Everything we know about LLMs.txt - Complete SEO, accessed December 21, 2025, <https://completeseo.com/everything-we-know-about-llms-txt/>
41. Instructor Adopts llms.txt: Making Documentation AI-Friendly, accessed December 21, 2025,  
<https://python.useinstructor.com/blog/2025/03/19/instructor-adopts-llms-txt/>
42. AI Crawler Access Checker - MRS Digital, accessed December 21, 2025,  
<https://mrs.digital/tools/ai-crawler-access-checker/>